

Step-by-step guide to Docker

Basic commands

1. Run the hello-world Docker container to verify basic functionality

```
$ docker run hello-world

Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
0e03bdcc26d7: Pull complete
Digest:
sha256:49a1c8800c94df04e9658809b006fd8a686cab8028d33cfba2cc049724254202
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
   (amd64)
3. The Docker daemon created a new container from that image which runs the
   executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it
   to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
```

2. Pull an image from Docker Hub

```
docker pull <image name>
```

EXAMPLE:

```
$ docker pull debian

Using default tag: latest
latest: Pulling from library/debian
e9afc4f90ab0: Pull complete
Digest:
sha256:46d659005ca1151087efa997f1039ae45a7bf7a2cbbbe2d17d3dcbda632a3ee9a
```

```
Status: Downloaded newer image for debian:latest
docker.io/library/debian:latest
```

3. Run a container and print OS information

```
docker run [options] <image name> <command>
```

We can display information about the OS by printing the `/etc/os-release` file:

```
$ docker run debian cat /etc/os-release

PRETTY_NAME="Debian GNU/Linux 10 (buster)"
NAME="Debian GNU/Linux"
VERSION_ID="10"
VERSION="10 (buster)"
VERSION_CODENAME=buster
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
```

Compare this information with those from your native OS.

4. Run an interactive shell inside a container

```
docker run -it <image name> bash
```

`-i` stands for interactive

`-t` allocates a pseudo-TTY

EXAMPLE:

```
$ docker run -it debian bash

root@9eed5b3d3044:/# whoami
root
root@9eed5b3d3044:/# ls
bin boot dev etc home lib lib64 media mnt opt proc root run sbin
srv sys tmp usr var
root@9eed5b3d3044:/# exit
```

5. List Docker images in the system

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
debian	latest	1b686a95ddb5	6 weeks ago	114MB
hello-world	latest	bf756fb1ae65	6 months ago	13.3kB

6. Run a container from an image with a tag different from `latest`

The general descriptor of Docker images is in the form `<user name>/<image name>:<image tag>`; in case of officially hosted images, the form is simply `<image name>:<tag>`. If the tag is not specified, Docker will default it to `latest`.

EXAMPLE:

```
$ docker pull debian:stretch

stretch: Pulling from library/debian
81fc19181915: Pull complete
Digest:
sha256:c6c98a905e230d779a92e6a329ff7ecad13b8ee0d21da8b013cee0df7e91c170
Status: Downloaded newer image for debian:stretch
docker.io/library/debian:stretch

$ docker run debian:stretch cat /etc/os-release

PRETTY_NAME="Debian GNU/Linux 9 (stretch)"
NAME="Debian GNU/Linux"
VERSION_ID="9"
VERSION="9 (stretch)"
VERSION_CODENAME=stretch
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
```

Compare this example with the one from point 2.

7. Write a simple Dockerfile

Dockerfiles are made of a sequence of commands to incrementally build the environment that will constitute a Docker image. They are similar to Bash scripts, with the addition of some specific keywords, called instructions. Every statement in a Dockerfile must start with an instruction.

The simplest and most useful instructions are:

- **FROM:** identify an already existing image as a base image; the subsequent instructions in the Dockerfile will add stuff on top of what's already defined in that image.
- **COPY:** copy files and directories from a source path into the image. The destination path will be automatically created if it does not exist.

- **RUN**: execute any command as if you were into a shell. RUN instructions usually make up the most of a Dockerfile, either installing software from the package manager or downloading and compiling resources.
- **ENV**: create a new environment variable in the image

EXAMPLE:

```
FROM debian:latest

RUN apt-get update && apt-get install -y wget

COPY script.sh /app

ENV NAME World
```

8. Build an image from a Dockerfile

```
docker build -t <name:tag> -f <Dockerfile path> <build context>
```

-t associates a user-supplied identifier to the new image. It is useful to already choose an identifier suitable for Docker Hub.

-f indicates the location of the Dockerfile to use. When the option is not provided, Docker defaults to a file called **Dockerfile** in the current directory. **-f** is thus useful when Dockerfiles have more elaborate names or reside in a different directory.

The build context is the set of files which will be available to the image builder. It usually corresponds to the current directory (i.e. **.**). The build context is used, for example, to find the files used by a **COPY** instruction.

EXAMPLE (with default Dockerfile and build context in the current directory):

```
$ docker build -t my_user/my_image:latest .
```

9. Login and push an image to Docker Hub

```
$ docker login
```

Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, head over to <https://hub.docker.com> to create one.
Username (<last logged user>):

Password:

Login Succeeded

```
$ docker push my_user/my_image:latest
```

Additional commands

1. List all Docker containers in the system (even stopped ones)

```
$ docker ps -a
```

CONTAINER ID	IMAGE	PORTS	COMMAND NAMES	CREATED
b8de0659bae5	debian:stretch		"cat /etc/os-release"	12 minutes ago
Exited (0) 12 minutes ago			hopeful_payne	
9eed5b3d3044	debian		"bash"	25 minutes ago
Exited (0) 25 minutes ago			zen_poitras	
dc0faa01b3a4	debian		"cat /etc/os-release"	26 minutes ago
Exited (0) 26 minutes ago			wizardly_herschel	
526cc2a156f6	hello-world		"/hello"	35 minutes ago
Exited (0) 35 minutes ago			admiring_mendeleev	

2. Run a container with automatic removal upon exit

By default, Docker does not delete containers after they complete the tasks assigned to them and return control to the shell. Instead, those containers remain in a stopped state, ready to be resumed if the user wishes so. To run a container that will be automatically removed when it exits, use the `--rm` option of `docker run`.

EXAMPLE:

```
$ docker ps -a
```

CONTAINER ID	IMAGE	PORTS	COMMAND NAMES	CREATED
b8de0659bae5	debian:stretch		"cat /etc/os-release"	12 minutes ago
Exited (0) 12 minutes ago			hopeful_payne	
9eed5b3d3044	debian		"bash"	25 minutes ago
Exited (0) 25 minutes ago			zen_poitras	
dc0faa01b3a4	debian		"cat /etc/os-release"	26 minutes ago
Exited (0) 26 minutes ago			wizardly_herschel	
526cc2a156f6	hello-world		"/hello"	35 minutes ago
Exited (0) 35 minutes ago			admiring_mendeleev	

```
$ docker run --rm debian:latest cat /etc/os-release
```

```
PRETTY_NAME="Debian GNU/Linux 10 (buster)"
NAME="Debian GNU/Linux"
VERSION_ID="10"
VERSION="10 (buster)"
VERSION_CODENAME=buster
ID=debian
HOME_URL="https://www.debian.org/"
SUPPORT_URL="https://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
```

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED
--------------	-------	---------	---------

STATUS	PORTS	NAMES	
b8de0659bae5	debian:stretch	"cat /etc/os-release"	15 minutes ago
Exited (0) 15 minutes ago		hopeful_payne	
9eed5b3d3044	debian	"bash"	28 minutes ago
Exited (0) 27 minutes ago		zen_poitras	
dc0faa01b3a4	debian	"cat /etc/os-release"	29 minutes ago
Exited (0) 29 minutes ago		wizardly_herschel	
526cc2a156f6	hello-world	"/hello"	38 minutes ago
Exited (0) 38 minutes ago		admiring_mendeleev	

3. Remove containers

```
docker rm <container ID or name> [<container ID or name>...]
```

EXAMPLE:

```
$ docker ps -a
```

CONTAINER ID	IMAGE	PORTS	COMMAND NAMES	CREATED
b8de0659bae5	debian:stretch		"cat /etc/os-release"	15 minutes ago
Exited (0) 15 minutes ago			hopeful_payne	
9eed5b3d3044	debian		"bash"	28 minutes ago
Exited (0) 27 minutes ago			zen_poitras	
dc0faa01b3a4	debian		"cat /etc/os-release"	29 minutes ago
Exited (0) 29 minutes ago			wizardly_herschel	
526cc2a156f6	hello-world		"/hello"	38 minutes ago
Exited (0) 38 minutes ago			admiring_mendeleev	

```
$ docker rm b8de0659bae5
```

b8de0659bae5

```
$ docker rm 9eed5b3d3044 dc0faa01b3a4 admiring_mendeleev
```

9eed5b3d3044
dc0faa01b3a4
admiring_mendeleev

```
$ docker ps -a
```

CONTAINER ID	IMAGE	PORTS	COMMAND NAMES	CREATED
STATUS				

A useful combination to remove **all** containers on the system with a single command is the following:

```
$ docker rm $(docker ps -aq)
```

4. Remove Docker images

```
docker rmi <image ID or name> [<image ID or name>...]
```

EXAMPLE:

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
debian	stretch	5738956efb6b	6 weeks ago	101MB
debian	latest	1b686a95ddbf	6 weeks ago	114MB
hello-world	latest	bf756fb1ae65	6 months ago	13.3kB

```
$ docker rmi 5738956efb6b 1b686a95ddbf hello-world
```

Untagged: debian:stretch
Untagged:
debian@sha256:c6c98a905e230d779a92e6a329ff7ecad13b8ee0d21da8b013cee0df7e91c170
Deleted:
sha256:5738956efb6b994ba2552e96100cc6a6d8a1f0c7dea402bddd54c5cf6f6871f8
Deleted:
sha256:8354d5896557e11ecd26e22b7845ecfed8fd6445c0362a0529466bf40f4b9a27
Untagged: debian:latest
Untagged:
debian@sha256:46d659005ca1151087efa997f1039ae45a7bf7a2cbbe2d17d3dcbda632a3ee9a
Deleted:
sha256:1b686a95ddbf7c369cbf099f60a1b75f0ca911466d8621f055bb4fdb3c6735b5
Deleted:
sha256:8803ef42039dcbe936755e9baae4bb7b19cb0fb6a438eb3992950cd0afef8e4f
Untagged: hello-world:latest
Untagged: hello-
world@sha256:49a1c8800c94df04e9658809b006fd8a686cab8028d33cfba2cc049724254202
Deleted:
sha256:bf756fb1ae65adf866bd8c456593cd24beb6a0a061dedf42b26a993176745f6b
Deleted:
sha256:9c27e219663c25e0f28493790cc0b88bc973ba3b1686355f221c38a36978ac63

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
------------	-----	----------	---------	------

5. Assign a different identifier to an existing image

```
docker tag <source image> <target image>
```

EXAMPLE:

```
$ docker tag dummy_image my_user/awesome_image:latest
```

Additional Dockerfile instructions

- **ADD:** copy files, directories and remote file URLs from a source into the image. It also automatically extracts tar archives into the image, which constitutes its best use case. Since the additional features of ADD are not immediately obvious, the official Docker documentation indicates `COPY` as the preferred instruction if files have to simply be transferred into an image.
- **WORKDIR:** set the working directory for subsequent instructions in the Dockerfile. If the WORKDIR doesn't exist, it will be created. Without a WORKDIR instruction, all actions in a Docker file happen at the filesystem root.
- **CMD:** provide default arguments for a container. These are the arguments used when nothing is entered after the image name in a `docker run` command. Providing arguments as part of `docker run` overrides the defaults set by `CMD`.
- **LABEL:** add metadata to an image in a key-value pair. An image can have multiple labels. Labels are additive, including labels in the base image indicated with FROM. Labels are useful for improved image classification and are sometimes used by third-party software.

Basic Dockerfile good practices

- **Do not use too many image layers:** Docker images are built from a series of layers, stacked on top of each other. Each layer represents an instruction in the image's Dockerfile and is simply a set of differences from the layer before it.

Try to achieve a balance between readability of the Dockerfile and reducing the number of image layers. Minimizing the number of layers also benefits total image size and performance of build and pull processes.

- **Cleanup after installations:** Because of the layered structure of Docker images, if files are downloaded and removed with different instructions, a copy of those files will still exist in the layer associated with the first instruction that retrieved them.

You can reduce the total image size by using a single RUN instruction that also cleans the package manager cache, or performs a complete installation from source and removes the original code.

Examples:

```
# Install from package manager and clean its cache
RUN apt-get update \
    && apt-get install -y --no-install-recommends \
        build-essential \
        wget \
    && rm -rf /var/lib/apt/lists/*

# Install from source and remove the code
RUN wget -q http://www.something.org/source-package.tar.gz \
    && tar xf source-package.tar.gz \
    && cd source-package \
    && ./configure \
    && make \
    && make install \
    && cd .. \
    && rm -rf source-package \
    && rm source-package.tar.gz
```


- **Avoid invalidating the build cache:** Each time `docker build` executes a Dockerfile instruction successfully, it caches the resulting image (even if it is an intermediate layer). When carrying out future builds of the Dockerfile, Docker will look for a match of a given instruction in its cache and, if found, it will reuse the cached layer instead of re-building it. Thus, proper use of the build cache can greatly speed up the creation of images.

Generally, if an instruction changes in a previously built Dockerfile, the lookup will fail and the build cache will be invalidated. When this happens, all subsequent Dockerfile commands will re-build new layers and the cache will not be used.

Special care should be used with `ADD` and `COPY` instructions: the contents of the files in the images are checksummed and, during cache lookup, the new checksum is compared against the checksum in the cached images. If anything has changed in the file(s), such as the contents and metadata, then the cache is invalidated. This means that even if the Dockerfile is identical, but you changed the files copied by an `ADD` or `COPY` instruction, a full rebuild will happen from that instruction onwards.

Building an "MPI Hello World" Docker image

In this extended example, we will package an "Hello World" MPI program in a Docker image. The program consists of a single C source file, called `hello_mpi.c` :

```
/* hello_mpi.c */

#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // Get the rank of the process
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Print off a hello world message
    printf("Hello world from rank %d of %d\n",
           rank, size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

It can be compiled with the accompanying Makefile, which uses the MPI compiler wrapper `mpicc` :

```
# Makefile

BIN=hello_mpi
MPICC?=mpicc

all: ${BIN}

hello_mpi: hello_mpi.c
    ${MPICC} -o hello_mpi hello_mpi.c

clean:
    rm ${BIN}
```

To create the container image, we need to write a Dockerfile covering the following points:

1. Provide a Linux distribution of choice
2. Install the necessary compilation tools
3. Install an MPI implementation
4. Assuming the "Hello MPI" program sources are available locally, copy them into the container
5. Call `make` to compile the "Hello MPI" program

A possible Dockerfile performing these tasks is provided below. The chosen Linux distribution is Debian 10 and the compilation toolchain is provided by the `build-essential` package. The MPI implementation of choice is MPICH 3.1.4, which is built from source after retrieving the corresponding archive with the `wget` utility.

```
# Dockerfile

FROM debian:buster

RUN apt-get update && apt-get install -y \
    build-essential \
    wget \
    ca-certificates \
    --no-install-recommends \
    && rm -rf /var/lib/apt/lists/*

RUN wget -q http://www.mpich.org/static/downloads/3.1.4/mpich-3.1.4.tar.gz \
    && tar xf mpich-3.1.4.tar.gz \
    && cd mpich-3.1.4 \
    && ./configure --disable-fortran --enable-fast=all,03 --prefix=/usr \
    && make -j$(nproc) \
    && make install \
    && ldconfig \
    && cd .. \
    && rm -rf mpich-3.1.4 \
    && rm mpich-3.1.4.tar.gz

COPY . /hello_mpi

RUN cd /hello_mpi && make
```

The image is then created with `docker build` :

```
docker build -t my_user/hello_mpi .
```

We can verify the correct functionality of the image by running multiple "Hello MPI" ranks with `mpirun` inside a container:

```
$ docker run --rm my_user/hello_mpi mpirun -n 4 /hello_mpi/hello_mpi  
  
Hello world from rank 0 of 4  
Hello world from rank 3 of 4  
Hello world from rank 2 of 4  
Hello world from rank 1 of 4
```